

# CSC3022H: Machine Learning

## Assignment 6

Reinforcement Learning

Department of Computer Science  
University of Cape Town

**Due: Friday, 29th May, 2020, 5.00 PM**

### Problem Description

Implement (in C++) the *Value Iteration* algorithm (detailed in chapter 3 [Sutton and Barto, 1998] and chapter 13 [Mitchell, 1997]) in order to find the optimal value ( $V^*$ ) for each state in a small grid-world (figure 1). Use the following information:

1. The agent has 4 actions  $\{ \textit{left}, \textit{right}, \textit{up}, \textit{down} \}$ , and the grid-world 6 states  $\{ s_1, s_2, s_3, s_4, s_5, s_6 \}$ . Figure 1 shows the possible transitions between states (actions for given states).
2. The state transition distribution  $P_{ss'}^a$  is deterministic, so  $P_{ss'}^a = 1.0$  for all states and actions.
3. Rewards for all state transitions are zero ( $R_{ss'}^a = 0$ ), except the following:

$$(1, 1) \rightarrow (2, 1); R_{ss'}^a = 50$$

$$(2, 0) \rightarrow (2, 1); R_{ss'}^a = 100$$

4. State  $s_3$  is the terminal state.
5. The discount factor is 0.8, **i.e.**  $\gamma = 0.8$ .




Figure 1: A small grid-world, where arrows show possible transitions between states. Note that state  $s_3$  is a terminal state.

**Question 1:** How many iterations does it take for the *Value Iteration* algorithm to converge? In an output text file list the optimal values ( $V^*$  for each state).

**Question 2:** Assume we start in state  $s_1$ , give the states that form the optimal policy ( $\pi^*$ ) to reach the terminal state ( $s_3$ ).

**Question 3:** Is it possible to change the reward function so that  $V^*$  changes, but the optimal policy ( $\pi^*$ ) remains unchanged?

If yes, describe how the reward function must be changed and the resulting change to  $V^*$ . Otherwise, briefly explain why this is impossible.

In a ZIP file, place the source code, makefile, and output text file (answers to questions 1, 2, 3).

Alongside your archive, please submit a file containing the hash, of your archive, produced by a MD5 checksum. See the link below for a tutorial on how to produce the hash:

<https://www.tecmint.com/generate-verify-check-files-md5-checksum-linux/>

Failure to include a '.md5' file waives your right to appeal your mark in relation to issues that arise from corrupted, missing or incorrect submissions.

## References

- [Mitchell, 1997] Mitchell, T. (1997). *Machine Learning: Chapter 13: Reinforcement Learning*. McGraw Hill, New York, USA.
- [Sutton and Barto, 1998] Sutton, R. and Barto, A. (1998). *An Introduction to Reinforcement Learning (Chapter 3)*. John Wiley and Sons, Cambridge, USA.